

Rule performance by metric — column-normalized (yellow=best, red=worst)
Annotated with raw median across all (base type $\times$ dataset $\times$ sample size) scenarios

